



RÉPUBLIQUE
FRANÇAISE

*Liberté
Égalité
Fraternité*

ONERA

THE FRENCH AEROSPACE LAB

Orchestrating Multi-Physical Simulations

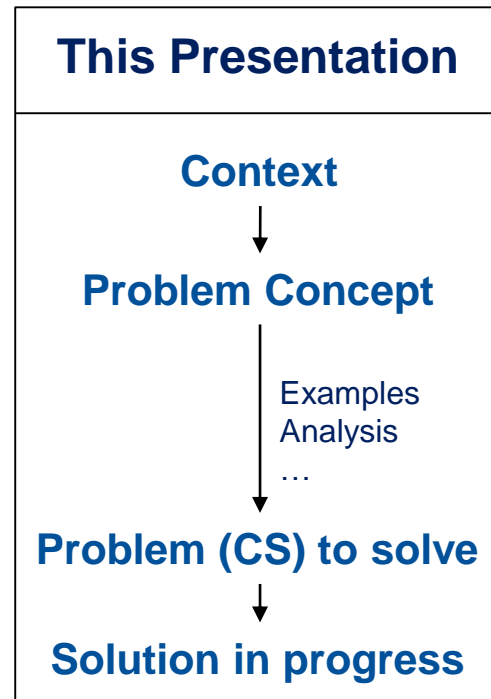
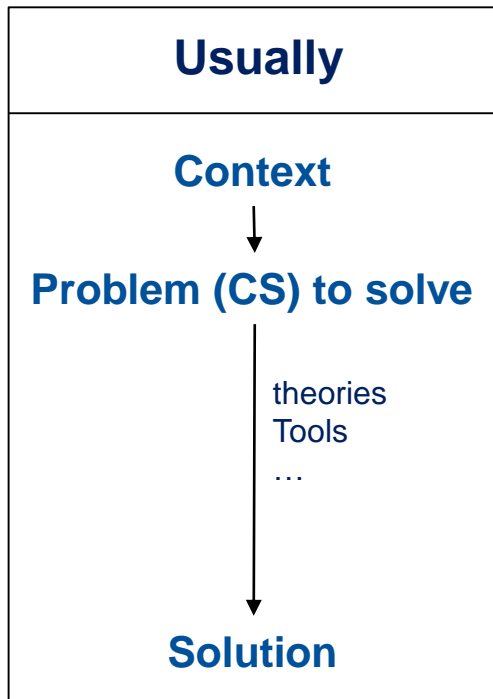
Michael Lienhardt,

ONERA, Université Paris-Saclay
France

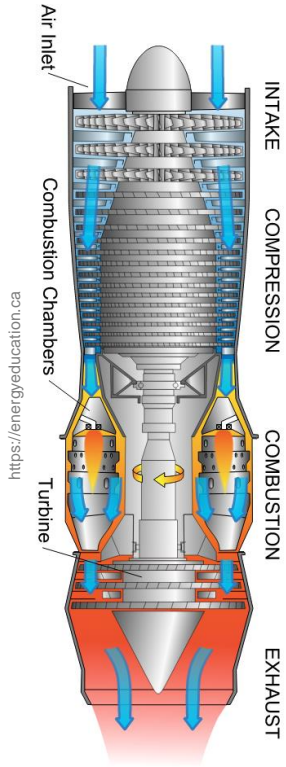
3rd October 2025

APM

Structure of this Presentation



Context: Physical Phenomena Simulation



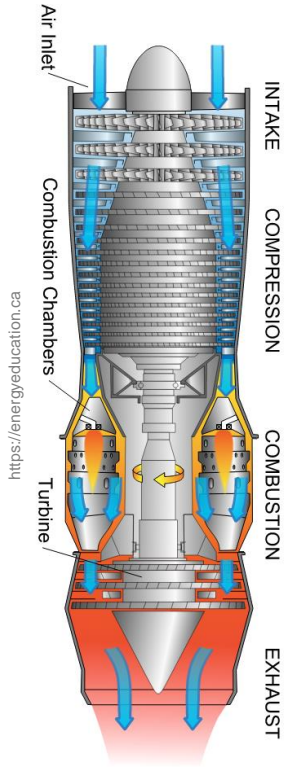
Principle of Simulation:

- Take the equations $f(t, x_i)=0$ encoding the different phenomena at play
- Discretize the space/time of simulation (e.g., using meshes and Δt)
- Implement the equations w.r.t. the discretization done in previous step
- Compute:
 - For all Δt to simulate:
 - Fixpoint to compute the x_i so that the $f(t, x_i)=0$ hold

Possibly many phenomena to take in account:

- Fluid Dynamics of different nature (air/gas/exhaust/rain/...)
- Solid (dilatation/corrosion/fatigue/...)
- Chemistry
- Energetic (e.g., to simulate lightnings or thermal signature)
- ...

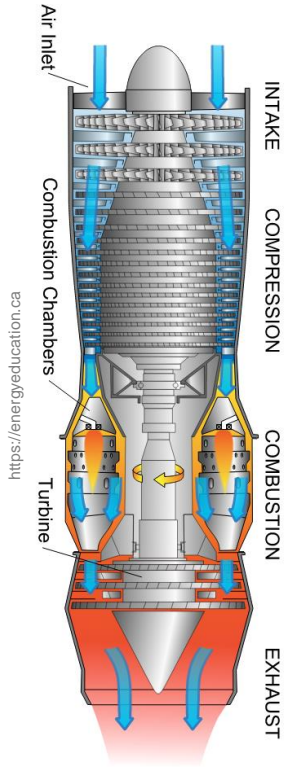
Context: Physical Phenomena Simulation



Problem of this approach: 1/3 - needs unrealistic computation space

- Each phenomena has its preferred space discretization
 - e.g., both fluid and energetic use mesh, but with different granularity
 - either too fine discretization,
or lots of transfer functions with approximation errors

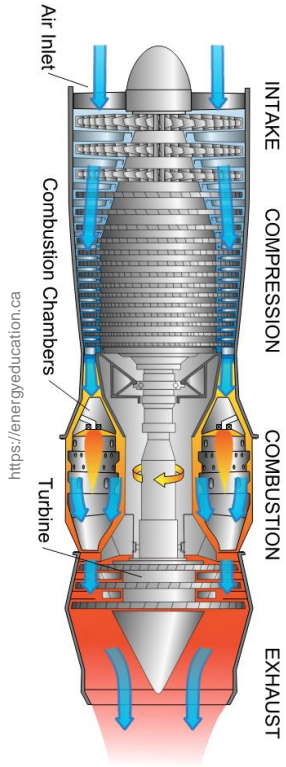
Context: Physical Phenomena Simulation



Problem of this approach: 2/3 - needs unrealistic computation time

- Each phenomena has its preferred time discretization
 - e.g., energetic is instantaneous, fluid is medium, solid is slow
 - using the greatest common divisor is too fine for most phenomena
 - also because Δt might change depending on the phenomena at play

Context: Physical Phenomena Simulation



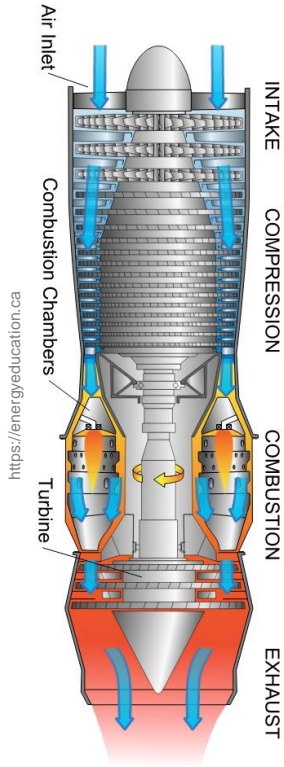
Problem of this approach:

3/3 - ad-hoc tricks specific to each phenomena

- Each phenomena has its specific tricks to make computation quicker
 - e.g., fluid do not simulate turbulence (very fine grain phenomena, chaotic), it generates it using functions tailored for the current simulation
- requires domain-specific knowledge and implementation

**In practice: each phenomena has its dedicated simulator...
and extra work needs to be done to simulate multiple phenomena**

Context: Physical Phenomena Simulation

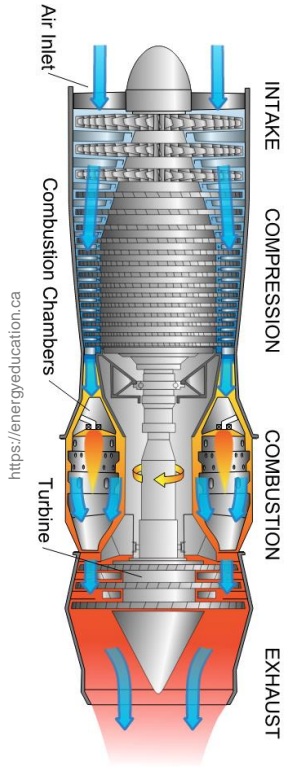


Multiple Phenomena Simulation: simulators orchestration

- Problem investigated for since the 70's
 - studied by apply mathematicians, physicists, numericians, ...
 - most solutions are hand-made and one-shot
 - helping tools: cwipi, Palm/OpenPalm
 - dead/live-locks are a real issue
- many computation problems, including
 - data interpolation between meshes of different shapes
 - techniques for different simulators to reach identical values
 - Δt adaptation heuristics to avoid simulators crashing

Context: Physical Phenomena Simulation

Multiple Phenomena Simulation: simulators orchestration

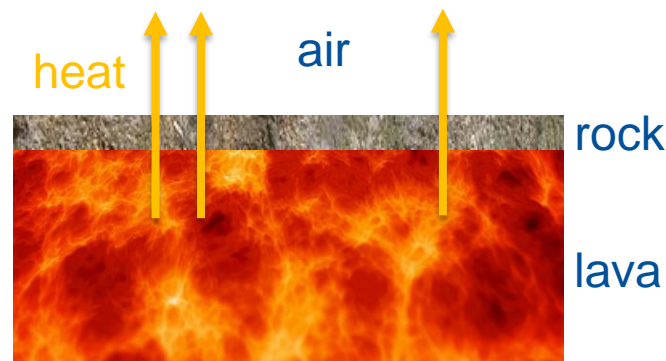


- Need to orchestrate Distributed Programs
 - some simulators are distributed, some not
 - orchestrator must be distributed to avoid bottle-neck
 - No known publication about the CS aspects of this orchestration

Problem Concept: Example (1/3)

Lava and Heat

- heat is transferred from the lava to the rock, and to the air
 - the opposite is technically true, but negligible
 - need for communications between the different simulators
-
- the Δt of every simulators are different
 - three simulators interact over two “interfaces”:
 - need to ensure that the computed values on both sides of each interface match



Problem Concept: Example (2/3)

Shard and Lightning

- Which wind pressure the Shard must deal with
- Δt is about 10ms

then, a lightning

- “instantaneously”, some gas becomes plasma
 - if nothing is done, the fluid simulator breaks:
 - too big discontinuities because too quick change
- Need to split its Δt and progressively insert the lightning



Problem Concept: Example (3/3)

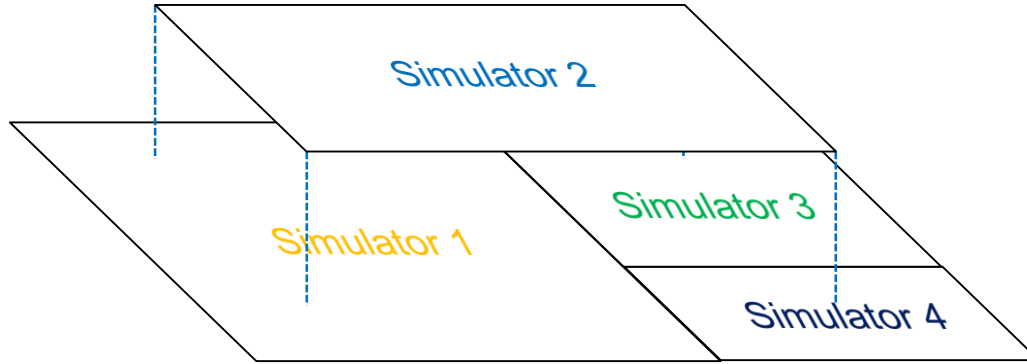
Autumn

- Suppose constant wind that pushes the leaves
- The wind pushes the leaves
- The leaves move and impact the air-flow



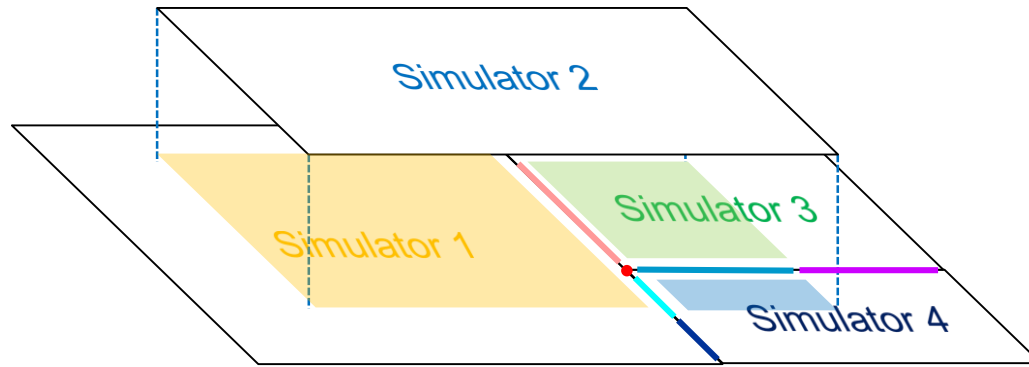
Problem Concept: Analysis

Simulators computation overlaps on the physical space



Problem Concept: Analysis

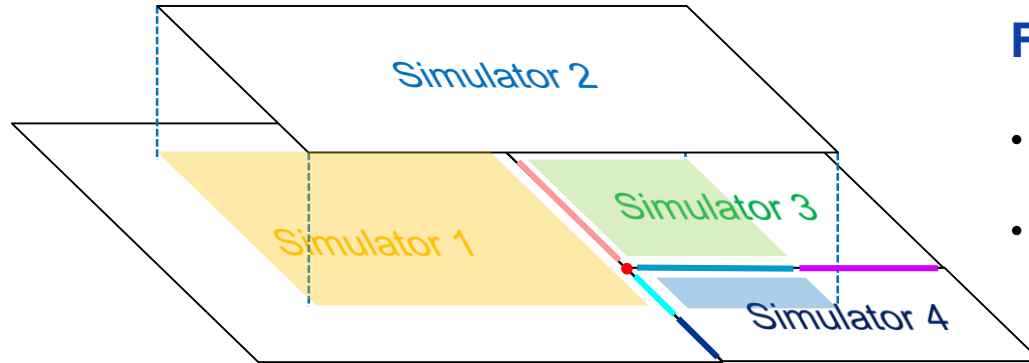
Simulators interact over “interfaces” (i.e., shared spaces)



Yellow	Volume: sim1 \Leftrightarrow sim2
Green	Volume: sim2 \Leftrightarrow sim3
Blue	Volume: sim3 \Leftrightarrow sim4
Red	Surface: sim1 \Leftrightarrow sim2 \Leftrightarrow sim3
Red	Point: sim1 \Leftrightarrow sim2 \Leftrightarrow sim3 \Leftrightarrow sim4
Cyan	Surface: sim1 \Leftrightarrow sim2 \Leftrightarrow sim4
Dark Blue	Surface: sim1 \Leftrightarrow sim4
Blue	Surface: sim2 \Leftrightarrow sim3 \Leftrightarrow sim4
Magenta	Surface: sim3 \Leftrightarrow sim4

Problem Concept: Analysis

Simulators interact over “Interfaces” (i.e., shared spaces)



For each Interface:

- Need a **consensus** over when to synchronize
- Need a **consensus** over computed values

(at every synchronization points)

Problem Concept: Analysis

Methods to reach a consensus: synchronization points

- Very physics dependent: how quickly every physics is changing the interface
- Related to the number of iterations each simulator need to reach a fixpoint
- Related to the CFL condition [1] that (partially) characterizes Δt
- only solution: fixpoint over reaching a solution
 - hypothesis: progress
 - hypothesis: fault-less network and agents

[1] https://en.wikipedia.org/wiki/Courant%E2%80%93Friedrichs%E2%80%93Lewy_condition

Problem Concept: Analysis

Methods to reach a consensus: computed values (1/2)

1. Use an Oracle that sets the initial values

- prescriptive: the values are the ones expected at the end of the computation
- indicative: the values should be close to the ones expected, but could be changed during computation
- Possible oracles:
 - none (uses the values of the previous computation step)
 - one of the simulator (e.g., in example 1 and 2)
 - ad-hoc functions

Problem Concept: Analysis

Methods to reach a consensus: computed values (2/2)

2. Use an ad-hoc correction function after computation

- collect the values computed by every solver, and sets common values
- prescriptive: the values are the ones expected at the end of the computation
- indicative: the values should be close to the ones expected, but maybe slightly off


3. Use an additional fixpoint

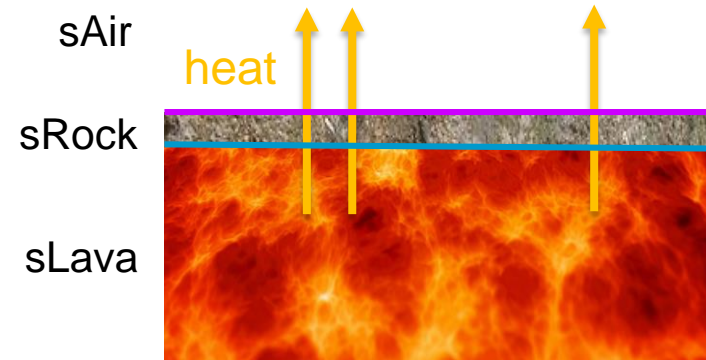
- the values produced after the correction is fed as input to the new iteration
 - these values “should” be closer to the expected ones
 - arguments for termination = theorems (in practice, should use a timeout, just in case)

Problem Concept: Analysis

Important details (1/5)



- Simulators as oracles = structural cause of deadlocks
 - Oracles must be executed before the simulators using their values
 - e.g., if specification error in example 1
 - in interface  :
 - Oracle(temperature) = sLava
 - Oracle(pressure) = sRock
- Deadlock

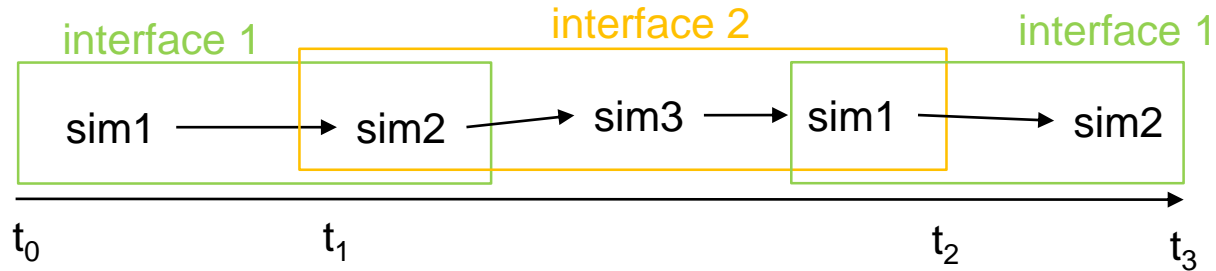


Problem Concept: Analysis

Important details (2/5)



- Simulators may start later, and end sooner
 - Would opposite dependencies in different interfaces also cause deadlocks?



Problem Concept: Analysis

Important details (3/5)



- Consensus can be “asymmetric”
- e.g., in example 2:
 - wind sets the speed of the leaves
 - leaves set the air-flow space



Problem Concept: Analysis

Important details (4/5)




- Fixpoint in interfaces can interact badly
 - Synchronization point = consensus over time advance

- e.g., in example 1:

- sRock and sAir agree to go to $t + \Delta t(\text{pink})$
- sLava and sRock cannot agree to go to $t + \Delta t(\text{blue})$
($\Delta t(\text{pink}) < \Delta t(\text{blue})$) and must restart computation on t

➤ “commit” inconsistency

➤ When sLava and sRock agree, next $\Delta t(\text{blue})$ must ensure progress (sRock is already at $t + \Delta t(\text{pink})$)



Problem Concept: Analysis

Important details (5/5)



- Failure management
 - Most of the time, failure are caused by simulators crashing/not converging
 - bugs
 - bad input values (i.e., the bug is silently from another simulator)
 - bad configuration (e.g., Δt too big or too small)
- Error recovery can be very difficult:
 - need to identify the cause of the problem (difficult)
 - need to roll-back to a safe state (data may be too large to accurately store a safe state)
 - or restart computation (which can take weeks) when the problem is identified

Problem to solve

- **Problem:** model the expected behavior of a simulation orchestration
- Abstraction:
 - Simulators, oracles, ad-hoc corrections = functions
 - in particular, the data updates performed by the correction in the simulators are abstracted away
 - Ordering between functions in an interface = DAG

Solution (in Progress)

Definition

An *OMPS* O is a tuple (F, D, I, L) where:

- F is a set of function names
- D gives the time interval $[f_b, f_e]$ for every function f in F
- I is a set of interfaces $i = (F_i, E_i)$ such that $F_i \subseteq F$ and i is a DAG
- L gives for every interface i the boolean i_l stating if it can loop

In the rest, we suppose given an OMPS $O = (F, D, I, L)$.

Solution (in Progress)

Definition

A *runtime state* is a pair $S = (S_f, S_i)$ where:

- S_f gives, for every function $f \in F$, the pair (f_o, f_u) of its last committed time and its current time
- S_i is a partial function which gives, for every interface $i \in I$, the pair (i_o, i_n) of its last committed time and its next committed time

Solution (in Progress)

Definition

A runtime state $S = (S_f, S_i)$ is *sound* iff:

- S_f gives consistent times for every functions, i.e., $\forall f \in F, f_o \leq f_u$
- S_i gives consistent times for every interface, i.e., $\forall i \in I, i_o \leq i_n$
- all interface must be considered:
 $\forall i \in I, \bigwedge_{f \in F_i} f_u \in [f_b, f_e[) \Leftrightarrow i \in \text{dom}(S_i)$
- the time of every function must be consistent w.r.t. the interfaces:
 $\forall i \in \text{dom}(S_i), \forall f \in F_i, i_o \leq f_o \wedge f_u \leq i_n$
- the DAG ordering must be respected:
 $\forall i \in \text{dom}(S_i), \forall (f, f') \in E_i, f_u < i_n \Rightarrow f'_u = i_o$

Solution (in Progress)

Definition

The initial runtime state of O is the pair (S_f, S_i) where:

- $S_f(f) = (f_b, f_b)$ for all $f \in F$
- $S_i = \emptyset$

Note: this initial state may not be sound...

Solution (in Progress)

Rule to start an “interface”

$$\begin{array}{c}
 \text{all functions are synchronized} \qquad \text{all functions can continue} \\
 \frac{
 \begin{array}{c}
 i \notin \text{dom}(S_i) \quad \overbrace{\forall f, f' \in F_i, f_u = f'_u} \\
 (S_f, S_i[i \mapsto (f_u, f_u)]) \triangleright_{f_u} (S'_f, S'_i) \quad \overbrace{\forall f \in F_i, f_u < f_e} \\
 \left. \vphantom{\frac{S_f, S_i[i \mapsto (f_u, f_u)] \triangleright_{f_u} (S'_f, S'_i)}} \right\} \text{Compute consensus}
 \end{array}
 }{
 (S_f, S_i) \triangleright (S'_f, S'_i)
 }
 \end{array}$$

Rule to consensus

$$\begin{array}{c}
 \text{all functions commit their time} \qquad \text{next rdv for interface given by consensus} \\
 (S_f, S_i) \triangleright_t (S_f[f \mapsto (f_u, f_u), f \in F_i], S_i[i \mapsto (i_n, \text{consensus}(t, F_i))])
 \end{array}$$

Solution (in Progress)

Rule to stop an “interface”

all functions are synchronized at interface's rdv

some function cannot continue

$$\frac{i \in \text{dom}(S_i) \quad \overbrace{\forall f, f' \in F_i, f_u = i_n} \quad \overbrace{\exists f \in F_i, f_u = f_e}}{\quad (S_f, S_i) \triangleright (S_f, S_i \setminus \{i\})}$$

Solution (in Progress)

Rule to continue an “interface”

$$\frac{\begin{array}{c} \text{all functions are synchronized at interface's rdv} \\ i \in \text{dom}(S_i) \quad \overbrace{\forall f, f' \in F_i, f_u = i_n} \\ \text{all functions can continue} \quad \overbrace{\forall f \in F_i, i_n < f_e} \\ \text{Compute consensus} \quad \overbrace{(S_f, S_i) \triangleright_{i_n} (S'_f, S'_i)} \end{array}}{(S_f, S_i) \triangleright (S'_f, S'_i)}$$

Problem to solve

Rule to reset an “interface”

$$\begin{array}{c}
 \text{can reset} \quad \text{all functions are synchronized} \quad \text{all functions can reset to last rdv} \\
 \text{at interface's rdv} \\
 \frac{
 \begin{array}{c}
 i \in \text{dom}(S_i) \quad L(i) \quad \forall f, f' \in F_i, f_u = i_n \quad \forall f \in F_i, f_o = i_o \\
 (S_f[f \mapsto (f_o, f_o), S_i[i \mapsto (i_o, i_o)]] \triangleright_{i_o} (S'_f, S'_i))
 \end{array}
 }{
 (S_f, S_i) \triangleright (S'_f, S'_i)
 }
 \end{array}
 \quad \text{Reset timers and compute consensus}$$

Problem to solve

Rule to stop an “interface”

$$\begin{array}{c}
 \text{must finish before its end} \quad \text{must finish before end of running interfaces} \quad \text{must finish before start of new interfaces} \\
 f \in F \quad u \in \left(\begin{array}{c} \underbrace{]f_u, f_e]}_{\text{must finish before its end}} \cap \underbrace{\bigcap_{i \in \text{dom}(S_i), f \in F_i}]f_u, i_n]}_{\text{must finish before end of running interfaces}} \cap \underbrace{\bigcap_{i \in I \setminus \text{dom}(S_i), f \in F_i}]f_u, f_b]}_{\text{must finish before start of new interfaces}} \end{array} \right) \\
 \underbrace{\bigwedge_{i \in \text{dom}(S_i), f \in F_i} (f', f) \in E_i \quad \bigwedge_{f'_u \geq i_n}}_{\text{Wait for dependencies}} \\
 \hline
 (S_f, S_i) \triangleright (S_f[f \mapsto (f_o, u)], S_i)
 \end{array}$$

Note: since this semantics put synchronization points at every function begin/end, opposite dependencies in different interfaces cause deadlocks

Conclusion

This Work:

- Studied the context of Orchestrating of Multiple Physical Simulations
 - From core principle of simulation, to illustrative examples
- Provided a model of interactions
 - Abstracting away data, communication and distribution
 - Focus on time and synchronization

Conclusion

Future Work:

- Study deadlock freedom
- Compare to existing implementations / tools
 - Refine the model
 - Suggest improvements
- Refine the notion of consensus (currently involves all nodes of all simulators)
- Consider recent developments, including
 - temporal interpolation
 - Simulators with local Δt

Thanks

- Jean-Didier Garaud
- Gilles Chaineray
- Karim Anemiche
- Eric Quémerais
- And many others...

Thank You

Questions?